

WRS: Waiting Room Sampling for Accurate Triangle Counting in Real Graph Streams

Dongjin Lee (dongjin.lee@kaist.ac.kr) and Kijung Shin (kijungs@kaist.ac.kr)

1 General Information

- Version: 2.0
- Date: Aug 18, 2020
- Authors: Dongjin Lee (dongjin.lee@kaist.ac.kr) and Kijung Shin (kijungs@kaist.ac.kr)

2 Introduction

WRS (Waiting Room Sampling) is a single-pass streaming algorithm for global and local triangle counting in (fully dynamic) real graph streams. **WRS** exploits a temporal dependency pattern in real dynamic graph streams. **WRS** has the following properties:

- *fast and any time*: WRS scales linearly with the number of edges in the input graph stream, and gives estimates at any time while the input graph grows
- *effective*: estimation error in WRS is up to 47% smaller than those in state-of-the-art methods
- *theoretically sound*: WRS gives unbiased estimates with small variance under the temporal locality.

Detailed information about the method is explained in the following paper

- Kijung Shin, “WRS: Waiting Room Sampling for Accurate Triangle Counting in Real Graph Streams”, IEEE International Conference on Data Mining (ICDM) 2017, New Orleans, USA
- Dongjin Lee, Kijung Shin, and Christos Faloutsos, “Temporal Locality-Aware Sampling for Accurate Triangle Counting in Real Graph Streams”, The VLDB Journal, 2020

3 Installation

- This package requires that java 1.7 or greater be installed in the system and set in PATH.
- For compilation (optional), type `./compile.sh`
- For packaging (optional), type `./package.sh`
- For demo (optional), type `make`

4 Input File Format for WRS_{INS} and WRS_{DEL}

4-1. Insertion-only Graph Stream

The input file lists edges in a graph in the order that they arrive. Each line corresponds to an edge and consists of a source-node id, a destination-node id, and a timestamp, which are separated by a tab. *example_graph.txt* is an example of the input file.

4-2. Fully Dynamic Graph Stream

The input file lists the edges in an *undirected* and *unweighted* graph in the order that they arrive. Each line corresponds to an edge addition or deletion. Each line consists of a source-node id, a destination-node id, and an indicator (1 for addition and -1 for deletion), which are integers separated by a tab. See *example_graph_dynamic.txt* for an example input file.

In both cases, we assume that there are **no parallel edges**. For example, both edge (1,2) and edge (2,1) cannot be in the input file at the same time. For a fully dynamic graph stream, if an edge has been added and has not been deleted yet, the same edge cannot be added.

5 Output Files Format for WRS_{INS} and WRS_{DEL}

Two output files are created.

- *global_count.out*: this file has the estimated number of global triangles.
- *Local_counts.out*: this file lists the estimated number of local triangles of each node. Each line consists of the node id and the number of its local triangle count, separated by a tab.

output_ins and *output_del* directory contains the examples of the output files.

6 Running WRS_{INS}

6.1 How to Run

```
./run_ins.sh input_path output_path k alpha
```

6.2 Parameters

- *input_path*: path of the input file. See 4 for the detailed format of the input file
- *output_path*: path of the directory for output files. See 5 for the detailed format of the output files
- *k*: maximum number of sampled edges (an integer greater than or equal to 2)
- *alpha*: the relative size of the waiting room (a real number in [0,1))

7 APIs for WRS_{INS}

7.1 Package: *wrs*

7.2 Class: *WRSIns*

7.3 Methods:

- `public WRSIns (int k, double alpha, int random_seed)`
 - create a *WRSIns* object
 - *k*: maximum number of sampled edges (an integer greater than or equal to 2)
 - *alpha*: the relative size of the waiting room (a real number in [0,1))
 - *random_seed*: an integer
- `public void processEdge (int src, int dst)`
 - process an edge
 - *src*: id of the source node
 - *dst*: id of the destination node
- `public double getGlobalTriangle ()`
 - return the estimated number of global triangles
- `public it.unimi.dsi.fastutil.ints.Int2DoubleMap getLocalTriangle ()`
 - return the estimated number of local triangles of each node
 - *return*: a map whose keys are node ids and values the estimated number of local triangle counts of the corresponding node.

7.4 Example Code: see *ExampleIns.java* for an example code using *WRSIns*.

8 Running WRS_{DEL}

8.1 How to Run

```
./run_del.sh input_path output_path k alpha
```

8.2 Parameters

- *input_path*: path of the input file. See 4 for the detailed format of the input file
- *output_path*: path of the directory for output files. See 5 for the detailed format of the output files
- *k*: maximum number of sampled edges (an integer greater than or equal to 2)
- *alpha*: the relative size of the waiting room (a real number in [0,1))

9 APIs for WRS_{DEL}

9.1 Package: *wrs*

9.2 Class: *WRSDel*

9.3 Methods:

- `public WRSDel (int k, double alpha, int random_seed)`
 - create a *WRSDel* object
 - *k*: maximum number of sampled edges (an integer greater than or equal to 2)
 - *alpha*: the relative size of the waiting room (a real number in [0,1))
 - *random_seed*: an integer
- `public void processEdge (int src, int dst, boolean add)`
 - process an edge
 - *src*: id of the source node
 - *dst*: id of the destination node
 - *add*: indicator (1 for addition and -1 for deletion)
- `public double getGlobalTriangle ()`
 - return the estimated number of global triangles
- `public it.unimi.dsi.fastutil.ints.Int2DoubleMap getLocalTriangle ()`
 - return the estimated number of local triangles of each node
 - *return*: a map whose keys are node ids and values the estimated number of local triangle counts of the corresponding node.

9.4 Example Code: see *ExampleDel.java* for an example code using *WRSDel*.